

## **Project Report: Image Classification with CIFAR-10**

### **1. Main Objective of the Analysis**

- **Objective:**
  - The primary goal of this project was to develop and evaluate deep learning models for classifying images from the CIFAR-10 dataset into one of ten predefined categories. To achieve this, I focused on utilizing Convolutional Neural Networks (CNNs), given their exceptional performance in image-related tasks. My analysis involved experimenting with different CNN architectures to enhance classification accuracy and understand their performance nuances.
- **Benefits:**
  - This analysis is crucial for advancing automated image recognition systems, which have applications across various sectors including autonomous vehicles, medical imaging, and security systems. By developing a robust image classification model, I aimed to deliver a system that not only performs accurately but also has the potential to generalize well to other datasets, thereby providing valuable insights and tools for the business and stakeholders.

### **2. Data Set Description**

- **Data Set Chosen:** CIFAR-10
- **Attributes:**
  - **Image Dimensions:** 32x32 pixels
  - **Color Channels:** 3 (RGB)
  - **Number of Classes:** 10
- **Class Labels:**
  - Airplane
  - Automobile
  - Bird
  - Cat
  - Deer
  - Dog
  - Frog
  - Horse
  - Ship
  - Truck
- **Objective:**
  - Train and evaluate CNN models for image classification.
  - Identify the best model architecture that balances accuracy and computational efficiency.

### **3. Data Exploration and Cleaning**

- **Exploration:**
  - **Data Distribution:** Analyzed the distribution of images across the 10 classes to ensure balanced representation.
  - **Visual Inspection:** Displayed sample images from each class to inspect the dataset's quality and diversity.
- **Cleaning and Preprocessing:**
  - **Normalization:** Scaled pixel values to [0, 1] by dividing by 255.
  - **Augmentation:** Applied techniques like random cropping, horizontal flipping, and rotation to increase training data diversity and reduce overfitting.
  - **Splitting:** Divided dataset into:
    - Training: 50,000 images
    - Validation: 5,000 images
    - Test: 5,000 images

## 4. Model Training

- **Model 1: Simple Convolutional Neural Network (CNN)**
  - **Architecture:**
    - **Conv2D Layers:**
      - Conv2D(32, (3, 3), activation='relu', input\_shape=(32, 32, 3))
      - Conv2D(64, (3, 3), activation='relu')
      - Conv2D(128, (3, 3), activation='relu')
    - **MaxPooling2D Layers:** MaxPooling2D((2, 2))
    - **Flatten Layer:** Flatten()
    - **Dense Layers:**
      - Dense(128, activation='relu')
      - Dropout(0.5)
      - Dense(10, activation='softmax')
  - **Hyperparameters:**
    - Learning Rate: 0.001
    - Batch Size: 256
    - Epochs: 10
  - **Performance:**
    - Accuracy: 69.32%
    - Loss: 0.8931
    - Macro avg Precision: 70.51%
    - Macro avg Recall: 69.32%
    - Macro avg F1-Score: 69.19%
    - Weighted avg Precision: 70.51%
    - Weighted avg Recall: 69.32%
    - Weighted avg F1-Score: 69.19%
- **Model 2: Deep Convolutional Neural Network (Deep CNN)**
  - **Architecture:**
    - **First Convolutional Block:**
      - Conv2D(64, (3, 3), activation='relu', padding='same')
      - BatchNormalization()

- Conv2D(64, (3, 3), activation='relu', padding='same')
  - BatchNormalization()
  - MaxPooling2D(2, 2)
  - Dropout(0.3)
  - **Second Convolutional Block:**
    - Conv2D(128, (3, 3), activation='relu', padding='same')
    - BatchNormalization()
    - Conv2D(128, (3, 3), activation='relu', padding='same')
    - BatchNormalization()
    - MaxPooling2D(2, 2)
    - Dropout(0.3)
  - **Third Convolutional Block:**
    - Conv2D(256, (3, 3), activation='relu', padding='same')
    - BatchNormalization()
    - Conv2D(256, (3, 3), activation='relu', padding='same')
    - BatchNormalization()
    - MaxPooling2D(2, 2)
    - Dropout(0.4)
  - **Fully Connected Layers:**
    - Flatten()
    - Dense(512, activation='relu')
    - Dropout(0.5)
  - **Output Layer:** Dense(10, activation='softmax')
- **Hyperparameters:**
  - Learning Rate: 0.001
  - Batch Size: 64
  - Epochs: 30
- **Performance:**
  - Accuracy: 73.27%
  - Precision: 0.7567
  - Macro avg Precision: 74.05%
  - Macro avg Recall: 73.27%
  - Macro avg F1-Score: 73.21%
  - Weighted avg Precision: 74.05%
  - Weighted avg Recall: 73.27%
  - Weighted avg F1-Score: 73.21%
- **Model 3: Residual Network (ResNet)**
  - **Architecture:**
    - **Input Layer:** Input(shape=x\_train[0].shape)
    - **First Convolutional Block:**
      - Conv2D(32, (3, 3), activation='relu', padding='same')
      - BatchNormalization()
      - Conv2D(32, (3, 3), activation='relu', padding='same')
      - BatchNormalization()
      - MaxPooling2D((2, 2))
    - **Second Convolutional Block:**

- Conv2D(64, (3, 3), activation='relu', padding='same')
- BatchNormalization()
- Conv2D(64, (3, 3), activation='relu', padding='same')
- BatchNormalization()
- MaxPooling2D((2, 2))
- **Third Convolutional Block:**
  - Conv2D(128, (3, 3), activation='relu', padding='same')
  - BatchNormalization()
  - Conv2D(128, (3, 3), activation='relu', padding='same')
  - BatchNormalization()
  - MaxPooling2D((2, 2))
- **Flattening Layer:** Flatten()
- **Fully Connected Layers:**
  - Dense(1024, activation='relu')
  - Dropout(0.2)
- **Output Layer:** Dense(10, activation='softmax')
- **Hyperparameters:**
  - Learning Rate: 0.0005
  - Batch Size: 128
  - Epochs: 50
- **Performance:**
  - Accuracy: 84.11%
  - Loss: 0.9241
  - Macro avg Precision: 84.30%
  - Macro avg Recall: 84.12%
  - Macro avg F1-Score: 84.11%
  - Weighted avg Precision: 84.30%
  - Weighted avg Recall: 84.12%
  - Weighted avg F1-Score: 84.11%

## 5. Recommended Model

- **Recommendation:** Residual Network (ResNet-50)
- **Rationale:**
  - ResNet's residual connections address the vanishing gradient problem in deeper networks, leading to better performance compared to simpler CNN and Deep CNN models.

## 6. Key Findings and Insights

- **Performance Improvement:**
  - Complex models with deeper architectures and residual connections significantly improve classification accuracy.
- **Class Confusion:**
  - Difficulty in distinguishing visually similar classes such as 'deer' and 'horse,' indicating areas for further refinement.

- **Implications:**
  - Need for additional fine-tuning or more sophisticated models, potentially incorporating advanced techniques like attention mechanisms or ensemble methods.

## 7. Suggestions for Next Steps

- **Model Enhancement:**
  - Experiment with advanced architectures such as DenseNet or EfficientNet.
- **Feature Engineering:**
  - Incorporate additional data augmentation techniques or use transfer learning with pre-trained models on larger datasets.
- **Error Analysis:**
  - Conduct detailed error analysis to understand misclassifications and refine the model or dataset.
- **Future Research:**
  - Explore semi-supervised or unsupervised learning techniques to leverage additional unlabeled data for further performance improvements.